

```
In [1]: !apt update -qq
!apt install nvidia-cuda-toolkit

buntu) focal/main amd64 libxkbfile1 amd64 1:1.1.0-1 [65.3 kB]
Get:14 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/main amd64 libxtst6 amd64 2:1.2.3-1 [12.8 kB]
Get:15 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/main amd64 libxxf86dga1 amd64 2:1.1.5-0ubuntu1 [12.0 k
B]
Get:16 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/main amd64 x11-utils amd64 7.7+5 [199 kB]
Get:17 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/main amd64 libatk-wrapper-java all 0.37.1-1 [53.0 kB]
Get:18 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/main amd64 libatk-wrapper-java-jni amd64 0.37.1-1 [45.1
kB]
Get:19 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/multiverse amd64 libcublaslt10 amd64 10.1.243-3 [9,249
kB]
Get:20 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/multiverse amd64 libcublas10 amd64 10.1.243-3 [29.7 MB]
Get:21 http://archive.ubuntu.com/ubuntu (http://archive.ubuntu.com/u
buntu) focal/multiverse amd64 libcudart10.1 amd64 10.1.243-3 [125 k
```

```
In [2]: !nvcc --version
```

```
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2022 NVIDIA Corporation
Built on Wed_Sep_21_10:33:58_PDT_2022
Cuda compilation tools, release 11.8, V11.8.89
Build cuda_11.8.r11.8/compiler.31833905_0
```

```

In [3]: %%writefile matrixMul.cu
#include <cmath>
#include <cstdlib>
#include <iostream>

#define checkCudaErrors(call)
do {
    cudaError_t err = call;
    if (err != cudaSuccess) {
        printf
("CUDA error at %s %d: %s\n", __FILE__, __LINE__,
cudaGetErrorString(err)); \
        exit(EXIT_FAILURE);
    }
} while (0)

using namespace std;

// Matrix multiplication Cuda
__global__ void matrixMultiplication
(int *a, int *b, int *c, int n) {
    int row = threadIdx.y + blockDim.y * blockIdx.y;
    int col = threadIdx.x + blockDim.x * blockIdx.x;
    int sum = 0;

    if (row < n && col < n)
        for (int j = 0; j < n; j++) {
            sum = sum + a[row * n + j] * b[j * n + col];
        }

    c[n * row + col] = sum;
}

int main() {
    int *a, *b, *c;
    int *a_dev, *b_dev, *c_dev;
    int n = 10;

    a = new int[n * n];
    b = new int[n * n];
    c = new int[n * n];
    int *d = new int[n * n];
    int size = n * n * sizeof(int);
    checkCudaErrors(cudaMalloc(&a_dev, size));
    checkCudaErrors(cudaMalloc(&b_dev, size));
    checkCudaErrors(cudaMalloc(&c_dev, size));

    // Array initialization
    for (int i = 0; i < n * n; i++) {
        a[i] = rand() % 10;
        b[i] = rand() % 10;
    }

    cout << "Given matrix A is =>\n";
    for (int row = 0; row < n; row++) {
        for (int col = 0; col < n; col++) {
            cout << a[row * n + col] << " ";
        }
        cout << "\n";
    }
    cout << "\n";
}

```

```

cout << "Given matrix B is =>\n";
for (int row = 0; row < n; row++) {
    for (int col = 0; col < n; col++) {
        cout << b[row * n + col] << " ";
    }
    cout << "\n";
}
cout << "\n";

cudaEvent_t start, end;

checkCudaErrors(cudaEventCreate(&start));
checkCudaErrors(cudaEventCreate(&end));

checkCudaErrors(cudaMemcpy
(a_dev, a, size, cudaMemcpyHostToDevice));
checkCudaErrors(cudaMemcpy
(b_dev, b, size, cudaMemcpyHostToDevice));

dim3 threadsPerBlock(n, n);
dim3 blocksPerGrid(1, 1);

// GPU Multiplication
checkCudaErrors(cudaEventRecord(start));
matrixMultiplication
<<<blocksPerGrid, threadsPerBlock>>>(a_dev, b_dev, c_dev, n);

checkCudaErrors(cudaEventRecord(end));
checkCudaErrors(cudaEventSynchronize(end));

float time = 0.0;
checkCudaErrors(cudaEventElapsedTime(&time, start, end));

checkCudaErrors(cudaMemcpy(c, c_dev, size, cudaMemcpyDeviceToHost));

// CPU matrix multiplication
int sum = 0;
for (int row = 0; row < n; row++) {
    for (int col = 0; col < n; col++) {
        sum = 0;
        for (int k = 0; k < n; k++) sum =
sum + a[row * n + k] * b[k * n + col];
        d[row * n + col] = sum;
    }
}

cout << "CPU product is =>\n";
for (int row = 0; row < n; row++) {
    for (int col = 0; col < n; col++) {
        cout << d[row * n + col] << " ";
    }
    cout << "\n";
}
cout << "\n";

cout << "GPU product is =>\n";
for (int row = 0; row < n; row++) {
    for (int col = 0; col < n; col++) {
        cout << c[row * n + col] << " ";
    }
}

```

```

        cout << "\n";
    }
    cout << "\n";

    int error = 0;
    int _c, _d;
    for (int row = 0; row < n; row++) {
        for (int col = 0; col < n; col++) {
            _c = c[row * n + col];
            _d = d[row * n + col];
            error += _c - _d;
            if (0 != (_c - _d)) {
                cout << "Error at
(" << row << ", " << col << ") => GPU:
" << _c << ", CPU: " << _d
                << "\n";
            }
        }
    }
    cout << "\n";

    cout << "Error : " << error;
    cout << "\nTime Elapsed: " << time;

    return 0;
}

```

Writing matrixMul.cu

```

In [4]: %%writefile matrixVectorMul.cu
#include <time.h>

#include <cmath>
#include <cstdlib>
#include <iostream>

#define checkCudaErrors(call)
    do {
        cudaError_t err = call;
        if (err != cudaSuccess) {
            printf
("CUDA error at %s %d: %s\n", __FILE__, __LINE__,
cudaGetErrorString(err)); \
            exit(EXIT_FAILURE);
        }
    } while (0)

using namespace std;

__global__ void matrixVectorMultiplication
(int *a, int *b, int *c, int n) {
    int row = threadIdx.x + blockDim.x * blockIdx.x;
    int sum = 0;

    if (row < n)
        for (int j = 0; j < n; j++) {
            sum = sum + a[row * n + j] * b[j];
        }

    c[row] = sum;
}

int main() {
    int *a, *b, *c;
    int *a_dev, *b_dev, *c_dev;
    int n = 10;

    a = new int[n * n];
    b = new int[n];
    c = new int[n];
    int *d = new int[n];
    int size = n * sizeof(int);
    checkCudaErrors(cudaMalloc(&a_dev, size * size));
    checkCudaErrors(cudaMalloc(&b_dev, size));
    checkCudaErrors(cudaMalloc(&c_dev, size));

    for (int i = 0; i < n; i++) {
        for (int j = 0; j < n; j++) {
            a[i * n + j] = rand() % 10;
        }
        b[i] = rand() % 10;
    }

    cout << "Given matrix is =>\n";
    for (int row = 0; row < n; row++) {
        for (int col = 0; col < n; col++) {
            cout << a[row * n + col] << " ";
        }
        cout << "\n";
    }
}

```

```

cout << "\n";

cout << "Given vector is =>\n";
for (int i = 0; i < n; i++) {
    cout << b[i] << ", ";
}
cout << "\n\n";

cudaEvent_t start, end;

checkCudaErrors(cudaEventCreate(&start));
checkCudaErrors(cudaEventCreate(&end));

checkCudaErrors(cudaMemcpy
(a_dev, a, size * size, cudaMemcpyHostToDevice));
checkCudaErrors(cudaMemcpy
(b_dev, b, size, cudaMemcpyHostToDevice));

dim3 threadsPerBlock(n, n);
dim3 blocksPerGrid(1, 1);

checkCudaErrors(cudaEventRecord(start));
matrixVectorMultiplication
<<<blocksPerGrid, threadsPerBlock>>>
(a_dev, b_dev, c_dev, n);

checkCudaErrors(cudaEventRecord(end));
checkCudaErrors(cudaEventSynchronize(end));

float time = 0.0;
checkCudaErrors(cudaEventElapsedTime(&time, start, end));

checkCudaErrors(cudaMemcpy(c, c_dev, size, cudaMemcpyDeviceToHost));

// CPU matrixVector multiplication
int sum = 0;
for (int row = 0; row < n; row++) {
    sum = 0;
    for (int col = 0; col < n; col++) {
        sum = sum + a[row * n + col] * b[col];
    }
    d[row] = sum;
}

cout << "CPU product is =>\n";
for (int i = 0; i < n; i++) {
    cout << d[i] << ", ";
}
cout << "\n\n";

cout << "GPU product is =>\n";
for (int i = 0; i < n; i++) {
    cout << c[i] << ", ";
}
cout << "\n\n";

int error = 0;
for (int i = 0; i < n; i++) {
    error += d[i] - c[i];
    if (0 != (d[i] - c[i])) {
        cout << "Error at (" << i << ") => GPU:

```

```
" << c[i] << ", CPU: " << d[i] << "\n";
    }
}

cout << "Error: " << error;
cout << "\nTime Elapsed: " << time;

return 0;
}
```

Writing matrixVectorMul.cu

```

In [5]: %%writefile vectorAdd.cu
#include <cstdlib>
#include <iostream>

#define checkCudaErrors(call)
do {
    cudaError_t err = call;
    if (err != cudaSuccess) {
        printf("CUDA error at %s %d: %s\n",
__FILE__, __LINE__, cudaGetErrorString(err)); \
        exit(EXIT_FAILURE);
    }
} while (0)

using namespace std;

// VectorAdd parallel function
__global__ void vectorAdd
(int *a, int *b, int *result, int n) {
    int tid = threadIdx.x + blockIdx.x * blockDim.x;
    if (tid < n) {
        result[tid] = a[tid] + b[tid];
    }
}

int main() {
    int *a, *b, *c;
    int *a_dev, *b_dev, *c_dev;
    int n = 1 << 4;

    a = new int[n];
    b = new int[n];
    c = new int[n];
    int *d = new int[n];
    int size = n * sizeof(int);
    checkCudaErrors(cudaMalloc(&a_dev, size));
    checkCudaErrors(cudaMalloc(&b_dev, size));
    checkCudaErrors(cudaMalloc(&c_dev, size));

    // Array initialization..You can use Randon function to assign values
    for (int i = 0; i < n; i++) {
        a[i] = rand() % 1000;
        b[i] = rand() % 1000;
        d[i] = a[i] + b[i]; // calculating serial addition
    }
    cout << "Given array A is =>\n";
    for (int i = 0; i < n; i++) {
        cout << a[i] << ", ";
    }
    cout << "\n\n";

    cout << "Given array B is =>\n";
    for (int i = 0; i < n; i++) {
        cout << b[i] << ", ";
    }
    cout << "\n\n";

    cudaEvent_t start, end;

    checkCudaErrors(cudaEventCreate(&start));
    checkCudaErrors(cudaEventCreate(&end));

```



```

    checkCudaErrors(cudaMemcpy
(a_dev, a, size, cudaMemcpyHostToDevice));
    checkCudaErrors(cudaMemcpy
(b_dev, b, size, cudaMemcpyHostToDevice));
    int threads = 1024;
    int blocks = (n + threads - 1) / threads;
    checkCudaErrors(cudaEventRecord(start));

    // Parallel addition program
    vectorAdd<<<blocks, threads>>>(a_dev, b_dev, c_dev, n);

    checkCudaErrors(cudaEventRecord(end));
    checkCudaErrors(cudaEventSynchronize(end));

    float time = 0.0;
    checkCudaErrors(cudaEventElapsedTime(&time, start, end));

    checkCudaErrors(cudaMemcpy
(c, c_dev, size, cudaMemcpyDeviceToHost));

    // Calculate the error term.

    cout << "CPU sum is =>\n";
    for (int i = 0; i < n; i++) {
        cout << d[i] << ", ";
    }
    cout << "\n\n";

    cout << "GPU sum is =>\n";
    for (int i = 0; i < n; i++) {
        cout << c[i] << ", ";
    }
    cout << "\n\n";

    int error = 0;
    for (int i = 0; i < n; i++) {
        error += d[i] - c[i];
        if (0 != (d[i] - c[i])) {
            cout << "Error at (" << i << ") =>
GPU: " << c[i] << ", CPU: " << d[i] << "\n";
        }
    }

    cout << "\nError : " << error;
    cout << "\nTime Elapsed: " << time;

    return 0;
}

```

Writing vectorAdd.cu

```
In [6]: !nvcc -dc matrixMul.cu
!nvcc *.o -o ./matrixMul && ./matrixMul
!rm -rf *.o
```

Given matrix A is =>

```
3 7 3 6 9 2 0 3 0 2
1 7 2 2 7 9 2 9 3 1
9 1 4 8 5 3 1 6 2 6
5 4 6 6 3 4 2 4 4 3
7 6 8 3 4 2 6 9 6 4
5 4 7 7 7 2 1 6 5 4
0 1 7 1 9 7 7 6 6 9
8 2 3 0 8 0 6 8 6 1
9 4 1 3 4 4 7 3 7 9
2 7 5 4 8 9 5 8 3 8
```

Given matrix B is =>

```
6 5 5 2 1 7 9 6 6 6
8 9 0 3 5 2 8 7 6 2
3 9 7 4 0 6 0 3 0 1
5 7 5 9 7 5 5 7 4 0
8 8 4 1 9 0 8 2 6 9
0 8 1 2 2 6 0 1 9 9
9 7 1 5 7 6 3 5 3 4
1 9 9 8 5 9 3 5 1 5
8 8 0 0 4 4 6 1 5 6
1 8 7 1 5 7 3 8 1 9
```

CPU product is =>

```
190 278 145 132 190 136 200 169 161 167
186 355 156 157 207 209 185 164 210 246
191 335 233 179 196 257 220 227 174 232
191 319 172 156 167 218 182 186 165 186
276 433 239 205 229 305 251 252 193 257
233 378 222 181 218 240 231 216 180 226
232 430 221 155 255 274 187 203 193 328
248 319 178 137 201 217 233 171 165 236
267 379 184 141 231 276 259 247 218 301
252 477 239 204 282 302 239 261 245 334
```

GPU product is =>

```
190 278 145 132 190 136 200 169 161 167
186 355 156 157 207 209 185 164 210 246
191 335 233 179 196 257 220 227 174 232
191 319 172 156 167 218 182 186 165 186
276 433 239 205 229 305 251 252 193 257
233 378 222 181 218 240 231 216 180 226
232 430 221 155 255 274 187 203 193 328
248 319 178 137 201 217 233 171 165 236
267 379 184 141 231 276 259 247 218 301
252 477 239 204 282 302 239 261 245 334
```

Error : 0

Time Elapsed: 0.026048

```
In [7]: !nvcc -dc matrixVectorMul.cu
!nvcc *.o -o ./matrixVectorMul && ./matrixVectorMul
!rm -rf *.o
```

```
Given matrix is =>
3 6 7 5 3 5 6 2 9 1
7 0 9 3 6 0 6 2 6 1
7 9 2 0 2 3 7 5 9 2
8 9 7 3 6 1 2 9 3 1
4 7 8 4 5 0 3 6 1 0
3 2 0 6 1 5 5 4 7 6
6 9 3 7 4 5 2 5 4 7
4 3 0 7 8 6 8 8 4 3
4 9 2 0 6 8 9 2 6 6
9 5 0 4 8 7 1 7 2 7
```

```
Given vector is =>
2, 8, 2, 9, 6, 5, 4, 1, 4, 2,
```

```
CPU product is =>
220, 147, 190, 201, 168, 171, 245, 235, 234, 210,
```

```
GPU product is =>
220, 147, 190, 201, 168, 171, 245, 235, 234, 210,
```

```
Error: 0
Time Elapsed: 0.017152
```

```
In [8]: !nvcc -dc vectorAdd.cu
!nvcc *.o -o ./vectorAdd && ./vectorAdd
!rm -rf *.o
```

```
Given array A is =>
383, 777, 793, 386, 649, 362, 690, 763, 540, 172, 211, 567, 782, 862, 6
7, 929,
```

```
Given array B is =>
886, 915, 335, 492, 421, 27, 59, 926, 426, 736, 368, 429, 530, 123, 13
5, 802,
```

```
CPU sum is =>
1269, 1692, 1128, 878, 1070, 389, 749, 1689, 966, 908, 579, 996, 1312,
985, 202, 1731,
```

```
GPU sum is =>
1269, 1692, 1128, 878, 1070, 389, 749, 1689, 966, 908, 579, 996, 1312,
985, 202, 1731,
```

```
Error : 0
Time Elapsed: 0.02048
```